

```
In [1]: #import Libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sb
import plotly.express as px
```

2. Load the dataset into the tool

```
In [3]: data = pd.read_csv('abalone.csv')
data
```

```
Out[3]:
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.1500	15
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.0700	7
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.2100	9
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.1550	10
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.0550	7
...
4172	F	0.565	0.450	0.165	0.8870	0.3700	0.2390	0.2490	11
4173	M	0.590	0.440	0.135	0.9660	0.4390	0.2145	0.2605	10
4174	M	0.600	0.475	0.205	1.1760	0.5255	0.2875	0.3080	9
4175	F	0.625	0.485	0.150	1.0945	0.5310	0.2610	0.2960	10
4176	M	0.710	0.555	0.195	1.9485	0.9455	0.3765	0.4950	12


4177 rows × 9 columns

#3. Perform Below Visualizations.

· Univariate Analysis

```
In [4]: data['Rings'].value_counts()
data.hist()
```


```
Out[4]: array([[<matplotlib.axes._subplots.AxesSubplot object at 0x7fde9ed90110>,
  <matplotlib.axes._subplots.AxesSubplot object at 0x7fde9ed522d0>,
  <matplotlib.axes._subplots.AxesSubplot object at 0x7fde9ed0a8d0>],
  [<matplotlib.axes._subplots.AxesSubplot object at 0x7fde9ecc0ed0>,
  <matplotlib.axes._subplots.AxesSubplot object at 0x7fde9ec84510>,
  <matplotlib.axes._subplots.AxesSubplot object at 0x7fde9ec43fd0>],
  [<matplotlib.axes._subplots.AxesSubplot object at 0x7fde9ebfe150>,
  <matplotlib.axes._subplots.AxesSubplot object at 0x7fde9ebb3750>,
  <matplotlib.axes._subplots.AxesSubplot object at 0x7fde9ebbfc50>]],
  dtype=object)
```



· Bi-Variate Analysis

```
In [5]: plt.scatter(data.Rings, data.Sex)
plt.title('The Gender of Abalone vs Number of Rings')
plt.xlabel('No. of Rings')
plt.ylabel('Gender')
```


```
Out[5]: Text(0, 0.5, 'Gender')
```



· Multi-Variate Analysis

```
In [6]: sb.heatmap(data.corr(),annot=True)
```

```
Out[6]: <matplotlib.axes._subplots.AxesSubplot at 0x7fde9e451b90>
```



#4. Perform descriptive statistics on the dataset.

```
In [7]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4177 entries, 0 to 4176
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype  
---  -
0   Sex              4177 non-null   object 
1   Length           4177 non-null   float64
2   Diameter         4177 non-null   float64
3   Height           4177 non-null   float64
4   Whole weight     4177 non-null   float64
5   Shucked weight   4177 non-null   float64
6   Viscera weight   4177 non-null   float64
7   Shell weight     4177 non-null   float64
8   Rings           4177 non-null   int64  
dtypes: float64(7), int64(1), object(1)
memory usage: 293.8+ KB
```

```
In [8]: data.describe()
```

```
Out[8]:
```

	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
count	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000
mean	0.523992	0.407881	0.139516	0.828742	0.359367	0.180594	0.238831	9.933684
std	0.120093	0.099240	0.041827	0.490389	0.221963	0.109614	0.139203	3.224169
min	0.075000	0.055000	0.000000	0.002000	0.001000	0.000500	0.001500	1.000000
25%	0.450000	0.350000	0.115000	0.441500	0.186000	0.093500	0.130000	8.000000
50%	0.545000	0.425000	0.140000	0.799500	0.336000	0.171000	0.234000	9.000000
75%	0.615000	0.480000	0.165000	1.153000	0.502000	0.253000	0.329000	11.000000
max	0.815000	0.650000	1.130000	2.825500	1.488000	0.760000	1.005000	29.000000

#5. Check for Missing values and deal with them.

There is no missing values

```
In [9]: data.isnull().any()
```

```
Out[9]: Sex                False
Length                False
Diameter              False
Height               False
Whole weight          False
Shucked weight        False
Viscera weight        False
Shell weight          False
Rings                 False
dtype: bool
```

#6. Find the outliers and replace them outliers

The dataset does not have a outliers

```
In [10]: fig = px.histogram(data, x='Whole weight')  
fig.show()
```

7. Check for Categorical columns and perform encoding.

There is one Categorical column SEX is replaced by an Integer

```
In [11]: from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
data["Sex"] = le.fit_transform(data["Sex"])
data["Sex"]
```

```
Out[11]: 0      2
1      2
2      0
3      2
4      1
..
4172   0
4173   2
4174   2
4175   0
4176   2
Name: Sex, Length: 4177, dtype: int64
```

#8. Split the data into dependent and independent variables.

```
In [12]: x=data.iloc[:,0:8].values
y=data.iloc[:,8:9].values
```

```
In [13]: x
```

```
Out[13]: array([[2.      , 0.455 , 0.365 , ..., 0.2245, 0.101 , 0.15  ],
 [2.      , 0.35  , 0.265 , ..., 0.0995, 0.0485, 0.07  ],
 [0.      , 0.53  , 0.42  , ..., 0.2565, 0.1415, 0.21  ],
 ...,
 [2.      , 0.6    , 0.475 , ..., 0.5255, 0.2875, 0.308 ],
 [0.      , 0.625 , 0.485 , ..., 0.531 , 0.261 , 0.296 ],
 [2.      , 0.71  , 0.555 , ..., 0.9455, 0.3765, 0.495 ]])
```



```
In [14]: y
```

```
Out[14]: array([[15],
               [ 7],
               [ 9],
               ...,
               [ 9],
               [10],
               [12]])
```

9. Scale the independent variables

```
In [15]: x=data.iloc[:,0:8]
print(x.head())
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	\
0	2	0.455	0.365	0.095	0.5140	0.2245	
1	2	0.350	0.265	0.090	0.2255	0.0995	
2	0	0.530	0.420	0.135	0.6770	0.2565	
3	2	0.440	0.365	0.125	0.5160	0.2155	
4	1	0.330	0.255	0.080	0.2050	0.0895	

	Viscera weight	Shell weight
0	0.1010	0.150
1	0.0485	0.070
2	0.1415	0.210
3	0.1140	0.155
4	0.0395	0.055

#10. Split the data into training and testing

```
In [16]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.3,random_state=0)
```

```
In [17]: x_train.shape
```

```
Out[17]: (2923, 8)
```

```
In [18]: x_test.shape
```

```
Out[18]: (1254, 8)
```

#11. Build the Model

```
In [19]: from sklearn.linear_model import LinearRegression  
lr = LinearRegression()
```

#12. Train the Model

```
In [20]: lr.fit(x_train, y_train)
```

```
Out[20]: LinearRegression()
```

#13. Test the Model

```
In [21]: y_pred = lr.predict(x_test)
print((y_test)[0:6])
print((y_pred)[0:6])
```

```
[[13]
 [ 8]
 [11]
 [ 5]
 [12]
 [11]]
[[13.11640829]
 [ 9.65691091]
 [10.35350972]
 [ 5.63648715]
 [10.67436485]
 [11.95341338]]
```

#14. Measure the performance using Metrics.

```
In [22]: # RMSE(Root Mean Square Error)

from sklearn.metrics import mean_squared_error
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
print("RMSE value : {:.2f}".format(rmse))
```

RMSE value : 2.26

```
In [23]: from sklearn.model_selection import cross_val_score
cv_scores = cross_val_score(lr, x, y, cv=5)
sco=cv_scores.round(4)
print(cv_scores.round(4))
print("Average",sco.sum()/5)
```

```
[0.4113 0.1574 0.4807 0.5046 0.4362]
Average 0.39803999999999995
```

